

# **Lean Consensus**

Thomas Coratger

# Contents

<b>1</b>	<b>Introduction: why consensus exists</b>	<b>1</b>
1.1	Blocks, chaining, and local verification . . . . .	1
1.2	What a consensus mechanism must guarantee . . . . .	2
1.2.1	Block creation and fork resolution . . . . .	2
1.2.2	Long-term guarantees . . . . .	2
1.2.3	From Proof of Work to Proof of Stake . . . . .	3
1.3	Lean Consensus as a redesign effort . . . . .	3
1.3.1	Ossification accelerationism . . . . .	4
1.4	How to read this document . . . . .	4
<b>2</b>	<b>The time model</b>	<b>5</b>
2.1	Slots . . . . .	5
2.2	Slot duration and intervals . . . . .	6
2.2.1	The four-interval workflow . . . . .	7
2.3	The slot clock . . . . .	8
2.3.1	Slot zero: the genesis slot . . . . .	8
2.4	Shorter slot rationale: The push for reduced latency . . . . .	8
2.5	Rationale: Explicit intervals and timing games . . . . .	10

## 1.0 Introduction: why consensus exists

Ethereum [1] is a distributed system: many independent computers (called nodes) communicate over a network that can be slow, unreliable, or adversarial. The purpose of a blockchain protocol is to let these nodes maintain a shared view of an evolving system state, even when some participants fail or misbehave.

A convenient mental model is a replicated state machine. A state machine is a function that takes (i) a current state and (ii) an input, and produces a new state. In Ethereum, the inputs are transactions (and related protocol messages), and the state contains things like account balances, smart contract code, and contract storage. The word replicated means that every honest node attempts to compute the same sequence of state transitions, so that they converge to the same result.

This raises an immediate challenge: in an open peer-to-peer network, different nodes may see different messages first, or may be targeted by conflicting information. Consensus is the part of the protocol whose job is to make the network agree on a single history of blocks. Once there is agreement on the block history, the order of transactions is fixed, and the state transitions become unambiguous.

### 1.1.0 Blocks, chaining, and local verification

The core data structure of a blockchain is a sequence of blocks. Each block bundles transactions together with metadata, and contains a cryptographic reference (a hash) to its parent block. This creates a chain: if any past block is modified, its hash changes, which breaks the link to all later blocks. Figure 1.1 illustrates this hash-linked structure.

Ethereum nodes do not accept blocks on trust. Instead, they verify blocks locally by re-executing the included transactions (and the protocol rules around them). This redundant verification is essential: it prevents a malicious peer from convincing you to accept an invalid state transition.

It is often useful to separate the protocol into two roles:

- A set of rules that define what it means for a block to be *valid* (the state transition function, plus any validity conditions).
- A set of rules that define which valid blocks are treated as *canonical* when there are competing histories (fork choice and finality).

In Ethereum today, these roles are commonly described as the execution layer and the consensus layer respectively. In this document we focus on the consensus side of the story, and in particular on a proposed redesign commonly referred to as Lean Consensus.

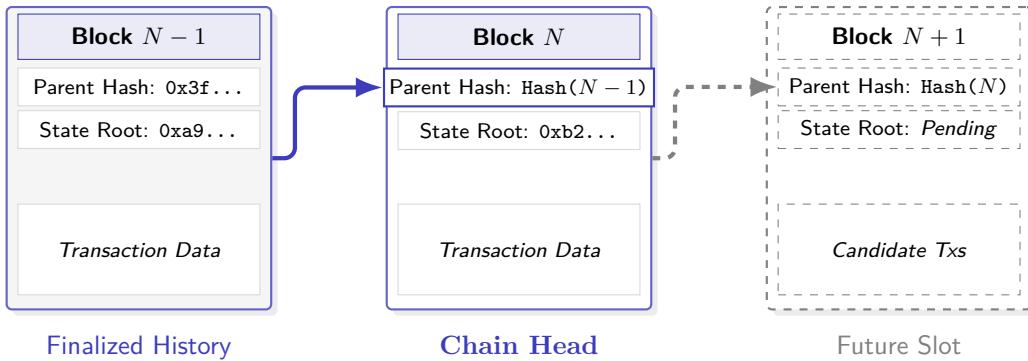


Figure 1.1: The cryptographic chain structure. Notice how the *Parent Hash* field of Block  $N$  explicitly commits to the entire content of Block  $N - 1$ . Changing any byte in Block  $N - 1$  would change its hash, breaking the link to Block  $N$ .

### 1.2.0 What a consensus mechanism must guarantee

In a distributed network, disagreement is not an exceptional event but a normal condition. Nodes do not receive messages at the same time, network links may be slow or unreliable, and different participants may temporarily observe different versions of the system state. If left unchecked, these differences would cause the network to fragment into incompatible histories.

A consensus mechanism exists to prevent this fragmentation from becoming permanent. Its role is not to eliminate disagreement entirely, but to ensure that honest nodes repeatedly reconverge on a single shared history of the blockchain.

#### 1.2.1 Block creation and fork resolution

To understand what this requires, it is helpful to separate the problem into two recurring situations that arise during normal operation.

1. *Block creation.* At any given moment, more than one participant may be capable of proposing a valid next block. The protocol must therefore specify who is allowed to propose blocks, when they may do so, and under which constraints. Without such rules, competing proposals would quickly overwhelm the system.
2. *Fork resolution.* Because messages propagate with delay, different nodes may temporarily build on different blocks. A consensus mechanism must specify how nodes choose between these competing histories so that the network eventually reconverges on a single chain. This process is illustrated in Figure 1.2.

#### 1.2.2 Long-term guarantees

Because a blockchain is intended to operate continuously, these mechanisms must satisfy long-term guarantees that hold over the lifetime of the system, not just in individual cases.

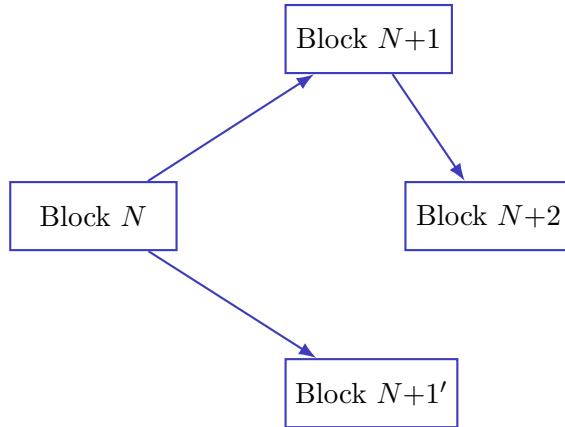


Figure 1.2: Temporary forks caused by network delays are resolved by the consensus rules, allowing the network to reconverge on a single chain.

1. *Safety*. Informally, safety means that the protocol should not lead honest nodes into an irreversible contradiction. Once part of the chain is considered final, honest nodes should not later finalize a conflicting history.
2. *Liveness*. Informally, liveness means that the system continues to make progress. Even in the presence of delays and partial failures, new valid blocks should keep being produced, and the chain should eventually advance.

There is one further requirement that is specific to public blockchains. Since anyone can join the network, the protocol must defend itself against Sybil attacks, where an adversary creates many fake identities to gain disproportionate influence. Consensus mechanisms address this by tying participation to a scarce resource. In Ethereum’s history, this resource was first external computation in Proof of Work, and later economic collateral in Proof of Stake.

### 1.2.3 From Proof of Work to Proof of Stake

In Proof of Work (PoW), the ability to propose blocks is tied to expending external resources (hardware and electricity). In Proof of Stake (PoS), it is tied to posting collateral inside the protocol itself.

Under PoS, the main actors are validators. A validator is a protocol participant that locks up ETH as stake and then runs software that proposes blocks and votes on blocks. If a validator violates key rules (for example, by equivocating—signing conflicting messages), the protocol can penalize them by slashing a portion of their stake. Figure 1.3 summarizes the shift in the resource that secures the system.

## 1.3.0 Lean Consensus as a redesign effort

Ethereum’s current PoS consensus layer is specified by the Beacon Chain design. Since that design was created, both the research landscape and the operational realities of Ethereum have evolved. A recurring motivation for Lean Consensus is to treat the consensus layer as something

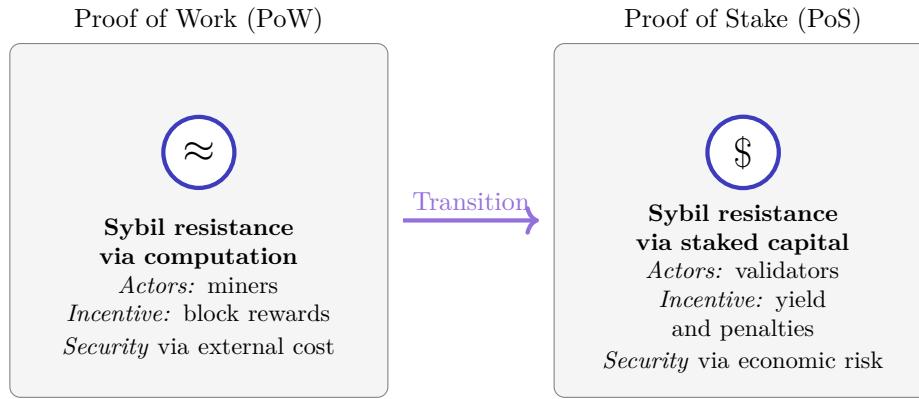


Figure 1.3: PoW and PoS use different Sybil-resistance resources. PoW relies on external expenditure; PoS relies on collateral that can be penalized by the protocol.

that can be simplified and modernized using lessons learned from the research front and real-world deployment.

Lean Consensus is commonly described as a holistic redesign. The intent is to revisit core components (such as validator responsibilities, committee structures, signatures, block production flow, etc) with a bias toward minimality, clearer invariants, and better alignment with the surrounding ecosystem.

### 1.3.1 Ossification accelerationism

Large protocol changes have coordination costs: they require client implementations, testing, audits, and social alignment. One proposed deployment philosophy for Lean Consensus is to batch major architectural changes into one big upgrade, instead of spreading them across many incremental hard forks.

The word ossification refers to reaching a stable maintenance mode where the core protocol changes rarely and only for well-understood reasons (such as security fixes). Ossification accelerationism is the argument that, if a major redesign is desirable, it can be better to do it deliberately in one concentrated effort, rather than living for many years with an increasingly complex intermediate design.

### 1.4.0 How to read this document

This text is written as an educational companion to the executable specifications in `leanSpec`. The goal is to make each concept self-contained before it is used. When new terminology is introduced, it will be defined and then reused consistently.

We proceed bottom-up: starting from the core abstractions (time, blocks, votes, finality, etc), then building toward the full Lean Consensus architecture and the rationale behind each design choice.

## 2.0 The time model

Distributed systems face a basic constraint: there is no global clock that all participants can read. Each node has its own local time, clocks drift, and messages arrive with unpredictable delays. As a result, questions like “which message came first?” do not have a universal answer at the network level.

Consensus protocols avoid relying on fine-grained timestamps. Instead, they introduce a discrete time model: the protocol defines a sequence of fixed-duration windows called *slots*. Protocol rules are expressed in terms of slot numbers, rather than in terms of real-time ordering. The only wall-clock quantity that nodes must share is a single anchor point, the *genesis time*, which marks the beginning of slot 0.

Lean Consensus uses a short slot duration and further subdivides each slot into smaller *intervals*. The purpose of this chapter is to make this time model precise: what slots and intervals are, how they are computed from wall-clock time, and why the protocol benefits from structuring time this way.

### 2.1.0 Slots

A *slot* is the protocol’s fundamental unit of time. Slots are numbered 0, 1, 2, … and advance deterministically. At any wall-clock time, every node can compute the current slot number locally from the genesis time and the slot duration. Figure 2.1 provides the basic picture: slots form a uniform grid laid over real time.

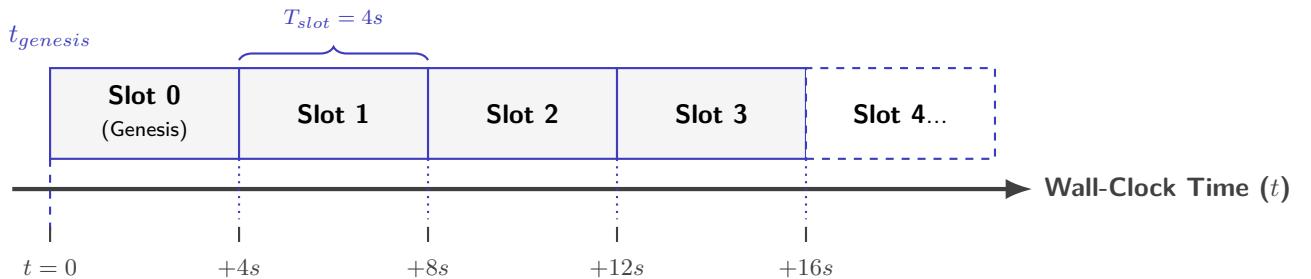


Figure 2.1: The relationship between physical wall-clock time and protocol slots. The genesis time  $t_{genesis}$  anchors the system. Every 4 seconds, the protocol deterministically transitions to the next slot number.

This discretization does not remove network delays, but it changes how they are handled. Instead of trying to totally order all events by timestamps, the protocol mostly cares about which slot an event belongs to. Events within the same slot are treated as belonging to the same time window,

which absorbs small timing differences and avoids turning clock skew into consensus ambiguity.

### 2.2.0 Slot duration and intervals

Lean Consensus fixes two public parameters:

- Slot duration:  $T_s = 4$  seconds.
- Interval duration:  $T_i = 1$  second.

Each slot is subdivided into  $T_s/T_i = 4$  intervals, numbered 0, 1, 2, 3. The interval structure gives the protocol a shared schedule for when particular messages are expected to be produced and when they are allowed to influence fork choice. Figure 2.2 shows one slot as four consecutive one-second phases.

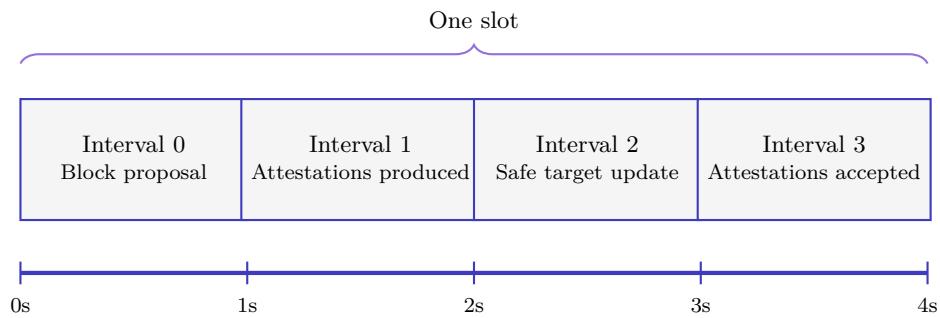


Figure 2.2: A slot is subdivided into four one-second intervals. Intervals define when different protocol actions occur.

The reason to stage actions like this is that time itself can become adversarial. In proof-of-stake protocols, validators can only attest to blocks they have actually received, and attestations must be produced before a fixed deadline within the slot. As a result, the moment at which a block is published determines how many validators see it in time to vote for it.

This dynamic is illustrated in Figure 2.3. A block proposer may benefit from delaying publication slightly, because waiting allows them to incorporate more information or extract more value. However, delaying also reduces the time available for the block to propagate through the network. The proposer is therefore trading off additional value against the risk that too few validators attest to the block. These strategic timing decisions, and their negative effects on consensus stability, are known in the Ethereum research literature as *timing games* [2].

A common mitigation pattern is therefore to make timing explicit and coarse-grained. Rather than letting messages influence consensus as soon as they arrive, the protocol defines clear phases within each slot and delays the point at which attestations begin affecting fork choice. This gives honest validators a predictable window to receive and verify blocks, and reduces the benefit of fine-grained message timing.




Figure 2.3: The race against the deadline. The vertical axis represents the percentage of the network that has received the block. An early proposal (blue) allows the gossip protocol to reach nearly all validators before the deadline. A late proposal (orange) leaves insufficient time for propagation, resulting in missed votes and potential consensus failure.

### 2.2.1 The four-interval workflow

The interval structure, as described in Figure 2.3, creates a pipeline that processes blocks and votes in a predictable sequence. This separation ensures that the network has time to reach a stable state before decisions are finalized.

1. **Interval 0: Block Proposal.** The slot begins with the publication of new data. The validator designated as the *proposer* creates a block and broadcasts it to the network.

Crucially, the proposer also casts their own vote (attestation) immediately. Instead of sending two separate messages—one for the block and one for the vote—the proposer packages their vote directly inside the block envelope (the data structure containing the block header and body). This optimization ensures the block comes with at least one vote of support attached.

2. **Interval 1: Attestation Broadcast.** Once the block is propagated, the remaining validators (the *attesters*) inspect it. If the block is valid, they broadcast their attestations.

At this stage, these incoming votes enter a pending state. The node stores them in memory but does not yet apply them to the consensus algorithm. Effectively, the node acknowledges receipt of the votes but does not yet allow them to influence the choice of the best chain.

3. **Interval 2: Safe Target Update.** Before counting the new votes, the protocol stabilizes its view of the chain history. It computes the safe target—the most recent block that has gathered a supermajority (at least 2/3) of the total stake.

The safe target acts as a high-confidence anchor. By identifying which block is "safe" before processing new votes, the protocol ensures that subsequent voting is built upon a stable foundation, reducing the likelihood of chaotic reorganizations.

4. **Interval 3: Attestation Acceptance.** Finally, the pipeline completes. The attestations that were held in the pending state during Interval 1 are promoted to known status.

The protocol now feeds these votes into the fork choice rule. The weights of the blocks are updated, and the node recomputes the canonical head—the tip of the chain with the most accumulated weight. This new head becomes the parent for the next slot’s proposal.

### 2.3.0 The slot clock

To make slot time concrete, the protocol needs a deterministic conversion from wall-clock time to a slot number and an interval number. Let:

- $t_g$  be the genesis time (a Unix timestamp, in seconds),
- $t$  be the current Unix time (in seconds),

The current slot is computed as:

$$\text{slot}(t) = \begin{cases} 0 & \text{if } t < t_g, \\ \left\lfloor \frac{t - t_g}{T_s} \right\rfloor & \text{otherwise.} \end{cases} \quad (2.1)$$

Within the current slot, define the offset into the slot:

$$\Delta(t) = (t - t_g) \bmod T_s, \quad (2.2)$$

and the corresponding interval:

$$\text{interval}(t) = \left\lfloor \frac{\Delta(t)}{T_i} \right\rfloor \in \{0, 1, 2, 3\}. \quad (2.3)$$

Figure 2.4 summarizes the flow from wall-clock time to protocol time.

### 2.3.1 Slot zero: the genesis slot

Slot zero holds a special place in the protocol. It is the *genesis slot*—the origin point of the chain. Unlike all subsequent slots, slot zero has no proposer and no parent block. The genesis state is created directly from configuration rather than derived from a previous state.

The genesis block serves as the trust anchor for the entire system. It is automatically considered both justified and finalized, bootstrapping the consensus mechanism. When the first real block arrives (at slot 1 or later), special logic acknowledges the genesis block as the immutable starting point from which all chain history flows.

### 2.4.0 Shorter slot rationale: The push for reduced latency

The decision to target a 4-second slot duration represents a significant acceleration from Ethereum’s historical 12-second heartbeat. While recent proposals such as EIP-7782 have explored conservative reductions to 6 or 8 seconds to improve throughput and user experience [3], Lean Consensus sets an aggressive engineering target of 4 seconds. This shift is motivated by three primary economic and structural factors identified in recent market microstructure research.

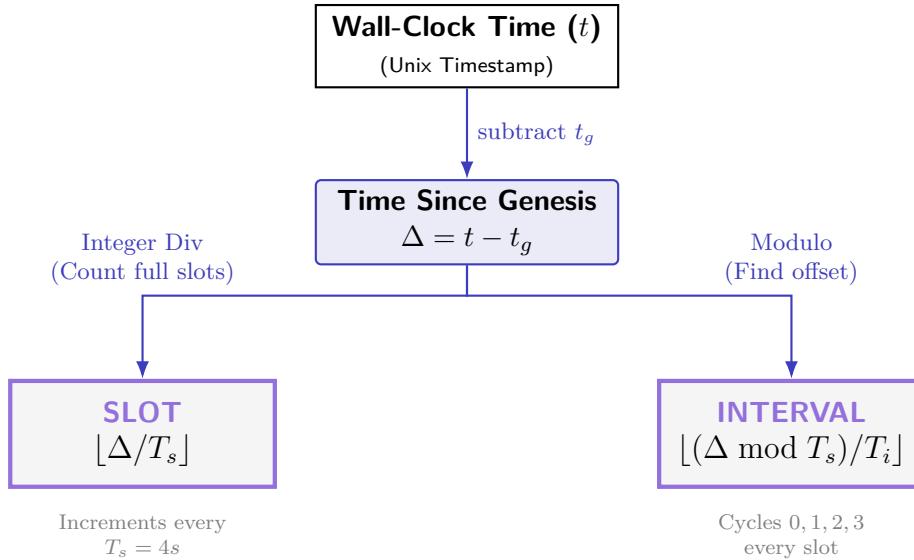


Figure 2.4: The slot clock logic. Wall-clock time is first normalized relative to genesis, then split into two coordinates: the Slot Number (coarse time) and the Interval (fine-grained offset within the slot).

- 1. The Settlement Layer (Confirmation Latency).** Ethereum’s primary role has evolved into a settlement layer for Layer 2 (L2) rollups. Rollups rely on the L1 chain to finalize their state or, at minimum, provide a high-confidence safe head—a block that has received a supermajority of validator votes and is thus statistically unlikely to be reorganized, even before full finality. The current 12-second latency forces cross-chain bridges and arbitrageurs to wait significant periods to ensure a transaction cannot be reorged. Reducing the slot time linearly reduces the time-to-inclusion and accelerates the generation of this safe confirmation signal that L2s read to derive the canonical state [4].
- 2. Economic efficiency (reducing price staleness).** On-chain automated market makers (AMMs) update their prices only when a new block is produced. Between blocks, prices remain fixed, even if the true market price of the asset moves elsewhere. By contrast, centralized exchanges update prices continuously as orders arrive.

This mismatch creates a predictable pattern of value extraction. When the market price moves between two blocks, arbitrageurs can trade against the AMM at the beginning of the next block, buying or selling at a price that is known to be outdated. These trades are risk-free for the arbitrageur and the resulting losses are borne by the AMM’s liquidity providers.

This effect is known as Loss-Versus-Rebalancing (LVR) [5]. Importantly, the magnitude of LVR depends on how long prices remain stale. Shorter slot times reduce the maximum price drift that can occur between blocks, shrinking the arbitrage opportunity available at each block boundary. Recent work shows that reducing this effect lowers the effective cost of providing liquidity on-chain, which can in turn support deeper and more efficient markets [6].

- 3. Resource Smoothing.** For a fixed transaction throughput (gas per second), shorter slots imply smaller block payloads. Propagating smaller payloads more frequently smooths out bandwidth spikes compared to propagating massive payloads infrequently. This reduction in peak bandwidth load helps maintain P2P network health and accessibility for nodes with constrained connections, even as total data throughput requirements grow [3].

### 2.5.0 Rationale: Explicit intervals and timing games

While shorter slots offer significant economic and utility benefits, they impose stricter requirements on the network. Specifically, reducing the slot time reduces the safety margin available for block propagation, making the consensus mechanism highly sensitive to *timing games*.

In a naive protocol, a block proposer is incentivized to delay their block publication as late as possible into the slot. By waiting, the proposer accumulates more order flow (MEV) and can capture more value. However, if a proposer broadcasts their block too late, honest validators may not receive it in time to verify and attest to it before the deadline. This behavior destabilizes the consensus layer, leading to missed slots and potential reorgs [2].

The rigid interval schedule described in Figure 2.2 acts as a counter-measure. By strictly enforcing a proposal deadline (Interval 0) separate from the attestation deadline (Interval 1), the protocol imposes a hard temporal boundary on block validity. If a proposer delays publication beyond the designated sub-slot to accumulate additional MEV, honest validators reject the block as untimely, thereby neutralizing the incentive for strategic delays and safeguarding the integrity of the 4-second cycle.

## Bibliography

- [1] V. Buterin *et al.*, “Ethereum white paper,” *GitHub repository*, vol. 1, no. 22-23, pp. 5–7, 2013.
- [2] casparschwa and mike, “Timing games: Implications and possible mitigations,” Ethereum Research (ethresear.ch), Dec. 2023, accessed: 2026-01-19. [Online]. Available: <https://ethresear.ch/t/timing-games-implications-and-possible-mitigations/17612>
- [3] B. Monnot, B. Adams, D. Feist, and J. Ma, “Eip-7782: The case for 2x shorter slot times in glamsterdam,” Ethereum Magicians Forum, Jun. 2025, accessed via provided context. [Online]. Available: <https://ethereum-magicians.org/t/eip-7782-the-case-for-2x-shorter-slot-times-in-glamsterdam/24616>
- [4] V. Buterin, “Epochs and slots all the way down: ways to give ethereum users faster transaction confirmation times,” Vitalik.ca, Jun. 2024. [Online]. Available: <https://vitalik.eth.limo/general/2024/06/30/epochslot.html>
- [5] J. Milionis, C. C. Moallemi, T. Roughgarden, and A. L. Zhang, “Automated market making and loss-versus-rebalancing,” *arXiv preprint arXiv:2208.06046*, 2022.
- [6] J. Ma and D. Crapis, “The cost of permissionless liquidity provision in automated market makers,” in *The International Conference on Mathematical Research for Blockchain Economy*. Springer, 2024, pp. 55–69.